

Closing the Security Activation Gap: A Security-by-Design Methodology for IoT Ecosystems and Its Empirical Validation in a DC Smart Microgrid

Raziel Campos-Sanchez^{a,b,*}, Luis Hernandez-Martinez^a, Claudia Feregrino-Uribe^{b,c}, Alejandro Penuelas-Angulo^{d,e}, Diego Cruz-Aguilar^{b,c}

^a*Electronics Department, National Institute of Astrophysics, Optics and Electronics (INAOE), Luis Enrique Erro 1, San Andrés Cholula, 72840, Puebla, Mexico*

^b*Science and Technologies of Security Department, INAOE, Luis Enrique Erro 1, San Andrés Cholula, 72840, Puebla, Mexico*

^c*Computational Sciences Department, INAOE, Luis Enrique Erro 1, San Andrés Cholula, 72840, Puebla, Mexico*

^d*CEIT-Basque Research and Technology Alliance (BRTA), Manuel Lardizabal 15, Donostia/San Sebastián, 20018, Guipuzcoa, Spain*

^e*Universidad de Navarra, Tecnum, Manuel Lardizabal 15, Donostia/San Sebastián, 20018, Guipuzcoa, Spain*

Abstract

A persistent security activation gap constrains IoT adoption: security mechanisms are widely available in commercial platforms yet remain inactive after standard deployment. We present a security-by-design methodology for centralized IoT ecosystems that treats identity as an intrinsic entity property, spans seven stages across four phases, and integrates cybersecurity throughout the development lifecycle. A dynamic Root of Trust process establishes verifiable cryptographic identity for each entity, providing the trusted foundation for authentication, access control, and secure communication.

We validate the methodology through a confirmatory case study applying all seven stages to a DC Smart Microgrid (SMG): two ESP32-based secondary control nodes and a Raspberry Pi 4B primary coordination agent.

*Corresponding author

Email addresses: rcampos@inaoe.mx (Raziel Campos-Sanchez), luish@inaoe.mx (Luis Hernandez-Martinez), cferegrino@inaoe.mx (Claudia Feregrino-Uribe), apenuelas@ceit.es (Alejandro Penuelas-Angulo), diego.cruz@inaoe.mx (Diego Cruz-Aguilar)

Four case study questions assess guidance sufficiency, identity model unambiguity, Root of Trust verifiability, and threat coverage. All seven stages produced documented, traceable artifacts with no external guidance required; a 15-parameter identity model across four entity types was produced without ambiguity; the methodology achieves unweighted OWASP IoT Top 10 coverage of 7/10 items fully addressed and 3/10 partially addressed. The Root of Trust, which reconstructs each node’s cryptographic identity from a Physically Unclonable Function response rather than a stored key, was validated end-to-end on real hardware for one node across $N = 30$ hard-reset trials (30/30 success, 558.0 ± 184.3 ms mean handshake latency); validating the second node remains future work. The methodology’s parameterized, technology-agnostic structure is designed to generalize across diverse IoT domains; empirical validation in additional domains remains future work.

Keywords: Internet of Things, Security-by-Design, IoT Methodology, Root of Trust, Identity Management, Smart Microgrid, Case Study Research

1. Introduction

The Internet of Things (IoT) has evolved from a conceptual vision into a technological paradigm underpinning critical infrastructure, industrial automation, healthcare systems, and smart cities (Al-Fuqaha et al., 2015; Zhou et al., 2021). Despite this expansion, large-scale IoT adoption remains constrained by fragmented architectures, heterogeneous technologies, and—critically—the failure to integrate cybersecurity from the earliest design stages (Shahinzadeh et al., 2024; Makhdoom et al., 2018).

The consequences of this gap are well documented. The Mirai botnet, which compromised over 600,000 IoT devices to launch unprecedented distributed denial-of-service attacks, originated from misconfigured or unprotected deployments relying on default credentials (Antonakakis et al., 2017). Subsequent analyses confirm that cyberattacks on IoT infrastructure continue to escalate in frequency and sophistication, targeting energy systems, healthcare networks, and public utilities (Abraham et al., 2025; Okika et al., 2025).

A central observation motivating this work is that the principal barrier to secure IoT deployment is not the *absence* of security features but the *security activation gap*: the disconnect between security capabilities *available* in a platform and those *operational* after a standard deployment (Figure 1).

Widely adopted open-source IoT platforms provide security mechanisms that require manual post-installation configuration, effectively restricting their use to organizations with dedicated security expertise (Fagan et al., 2021).

A second, more fundamental barrier operates at the methodological level: most IoT development projects lack a structured process that integrates cybersecurity from organizational decision-making through operational maintenance. Security is typically addressed *after* architectural and technology decisions are made, producing mechanisms that are bolt-on rather than foundational, with identity management that is platform-specific rather than entity-centric.

To address both barriers, we present two complementary contributions:

1. **A security-by-design methodology** for developing centralized IoT ecosystems. The methodology treats identity as an intrinsic property of every entity, spans seven stages across four phases (organizational strategy, architecture, implementation, and operation), and embeds cybersecurity—from threat policy to Root of Trust establishment—as a foundational concern from problem definition through ongoing operation. A dynamic Root of Trust (RoT) process establishes verifiable cryptographic identity for each entity before any operational communication begins.
2. **An empirical confirmatory case study** applying the methodology to the design and implementation of a DC Smart Microgrid (SMG) IoT ecosystem, following Yin’s (2014) Case Study Research design. The SMG comprises two ESP32 secondary control nodes and a Raspberry Pi 4B primary coordination agent, reaching Technology Readiness Level 3. Four case study questions evaluate guidance sufficiency, identity model unambiguity, Root of Trust verifiability, and OWASP IoT threat coverage.

Conceptual foundation: The identity-centric IoT paradigm, the IoT/ecosystem distinction, the formal entity identity model, and the reconceptualization of the Root of Trust as a multi-phase process are established at the conceptual and definitional level in a companion manuscript (Campos-Sanchez et al., 2026a). The present work builds directly on that foundation to specify, validate, and empirically evaluate an operative seven-stage development methodology; it does not re-derive these definitions independently.




Figure 1: The security activation gap: standard IoT deployments activate only a subset of platform security controls by default; the remaining controls require manual post-installation configuration by practitioners with dedicated security expertise. The proposed methodology activates all controls by design, as a direct outcome of following the seven development stages.

Scope note: The ESP32/C (ESP-IDF) firmware suite and the primary agent software implementing the SMG are presented as part of this work’s own case study rather than as a companion contribution under separate review. The present work evaluates the *methodology* through the case study; implementation details are provided only to the extent required for methodological evidence.

We evaluate the methodology through: (a) stage-by-stage traceability, demonstrating that all seven stages produce documented, actionable artifacts; (b) quantitative measurements of identity provisioning, handshake latency, sampling performance, and heap stability; and (c) a systematic OWASP (Open Web Application Security Project) IoT Top 10 coverage analysis.

The remainder of this article is organized as follows. Section 2 reviews related work. Section 3 presents the conceptual framework. Section 4 details the methodology and its STRIDE threat coverage model. Section 5 presents the SMG case study. Section 6 reports empirical results. Section 7 discusses findings, and Section 8 concludes.

2. Related Work

We adopt the definition of IoT methodology from Fortino et al. (2020): a structured set of guidelines, methods, practices, and documents organized within a logical framework for solving the specific problem of IoT ecosystem development. We evaluate existing methodologies against nine criteria: the eight classical dimensions drawn from the literature, supplemented by a ninth, *security validation*, which captures whether the methodology’s security provisions are supported by empirical evidence from a formal case study or experiment.

Table 1 compares representative methodologies. A 2023 comparative study (Moreta et al., 2023) confirms cybersecurity integration remains the most underdeveloped dimension across IoT development methodologies: most proposals emphasize functional aspects—interoperability, scalability, technological integration—while treating cybersecurity as an afterthought or omitting it entirely (Fortino et al., 2020; Al-Fuqaha et al., 2015; Merzouk et al., 2020). Agent-based approaches (Manate et al., 2014; Ayala et al., 2015; Spanoudakis and Moraitis, 2015) provide structured development processes but lack explicit security mechanisms and identity assignment; model-driven methodologies (Morin et al., 2017; Costa et al., 2016; Ciciirelli et al., 2017)

address code heterogeneity and portability but not entity identity or security. A systematic review of agile IoT methodologies ([Guerrero-Ulloa et al., 2023](#)) confirms even recent structured approaches defer authentication, identity management, and threat modeling until after functional requirements are met—if addressed at all.

Security-focused proposals remain limited in scope and depth. [Josyula and Gupta \(2017\)](#) presents a security methodology for IoT systems engineering without specific guidelines for identity assignment or cryptographic mechanism selection. [Savaglio et al. \(2018\)](#) includes security aspects (authorization, key management, authentication) but not selection criteria or their timing relative to development stages. [Muñoz et al. \(2025\)](#) proposes a five-layer service-oriented architecture with security as a cross-cutting concern, without explaining how mechanisms are integrated or selected; its PKI-based scheme neglects resource-constrained devices. A systematic review of IoT authentication schemes in this journal ([Saqib and Moon, 2023](#)) (over 120 proposals, 2010–2022) confirms core fragmentation: authentication mechanisms are proposed in isolation, without connection to a broader development process, formal identity model, or measurable threat coverage metric. [Blanco et al. \(2023\)](#) propose an ontology-driven framework for cyber-physical security requirements, providing structured elicitation but not a full lifecycle spanning organizational policy through operational deployment.

More recently, several efforts target specific IoT security domains without holistic development guidance: Zero Trust architectures for industrial IoT ([Zanasi et al., 2024](#); [Fernandez and Brazhuk, 2024](#)), blockchain-based supply chain security ([Ray et al., 2024](#); [Fernando et al., 2023](#)), and SDN-based IoT fog security ([Javanmardi et al., 2023](#); [Turner et al., 2023](#)) address individual threat surfaces but not a development methodology. [Visintin \(2024\)](#) examines identity of things device-centrally but does not extend to ecosystem-wide identity management across heterogeneous entity types or connect identity to access control policy. PUF-based Root of Trust implementations ([Sassalou et al., 2025](#)) and hardware security modules ([Bakiris and Oikonomakis, 2025](#)) strengthen device-level identity but remain single-layer solutions, leaving the broader development process unaddressed. This work adopts an SRAM-based PUF for the ESP32, `esp32_puflib` ([Stanicek, 2022](#)), as the concrete Stage 4 technology for identity-parameter generation without persistent on-device key storage (Section 4), integrated into the full seven-stage process rather than treated as an isolated device-level mechanism.

Standards from government agencies and consortia—Secure-by-Design (UK Government Security, 2024), NIST IoT Cybersecurity Guidance (Fagan et al., 2021), ETSI EN 303 645 (European Telecommunications Standards Institute, 2020), the CSA IoT Security Controls Framework (Cloud Security Alliance, 2023), and the 2025 NIST cybersecurity supply chain risk guidance (Licata et al., 2025)—provide valuable security baselines. However, these standards remain predominantly prescriptive at the policy level, offering no technology selection criteria, no process for deriving implementation choices from organizational constraints, and no deployment automation.

Taken together, these approaches share a common limitation: the absence of a methodology that simultaneously addresses (a) organizational adoption strategy, (b) technology-agnostic architectural design with a formal entity-centric identity model, (c) a dynamic Root of Trust process establishing verifiable cryptographic identity, (d) iterative development guided by maturity levels, and (e) empirical validation through a formal case study with measurable artifacts at every stage. The methodology presented in this work addresses all five dimensions. The validation follows Yin’s (2014) confirmatory case study design with four research questions and documented stage-by-stage artifact traceability—we are not aware of a comparable empirical validation of a complete IoT security development methodology using this research design.

3. Conceptual Framework

Ambiguity in IoT definitions hinders coherent methodology design (Sorri et al., 2022). Rather than re-deriving reference definitions, this methodology builds directly on the identity-centric conceptual framework established in a companion manuscript (Campos-Sanchez et al., 2026a), which reframes IoT as a technological paradigm in which cyber-physical systems are constitutively endowed with identity, and distinguishes this paradigm from an *IoT ecosystem*—a digital environment in which entities possessing unique identities interact under a defined structure to provide services that solve identified problems. We adopt both definitions unchanged and use them as the entry point for the present methodology.

The ecosystem comprises seven element types: IoT devices (embedded systems with ecosystem-relative identity), IoT servers (network-connected infrastructure enabling entity communication), IoT services (processes enabling ecosystem operation and security), storage systems (information persistence

Table 1: Comparison of IoT ecosystem development methodologies (9 criteria). Tech.Dev. = technological development guidance; CyberSec. = explicit cybersecurity integration; Agnostic = technology-agnostic design; Org. = organizational aspects; Scalability; Maturity = TRL or equivalent; Interop. = interoperability; Concept. = formal conceptual framework; Sec.Val. = security provisions validated through formal case study or experiment. ✓ full, ~ partial, ✗ absent. *This work* is rated ~ on Tech.Dev. because the methodology is technology-agnostic by design: it specifies mechanism categories and stage requirements, not specific technology selection; ~ on Maturity and Scalability because the case study reaches only TRL 3 (laboratory proof of concept) on 2 nodes, not a production-scale or externally-validated deployment. Ratings for “This work” are self-assessed by the authors, not externally validated.

Methodology	Tech. Dev.	Cyber Sec.	Agnostic	Org.	Scalability	Maturity	Interop.	Concept. Frame.	Sec. Val.
Manate et al. (2014) (Manate et al., 2014)	—	—	✗	✗	✓	✗	✓	✗	✗
Ayala et al. (2015) (Ayala et al., 2015)	✓	✗	✓	✗	—	✗	✓	—	✗
Spanoudakis et al. (2015) (Spanoudakis and Moraitis, 2015)	✓	✗	✗	—	✓	✗	✓	✗	✗
Patel et al. (2015) (Patel and Cassou, 2015)	✓	✗	✗	✗	✓	✗	✓	✓	✗
Costa et al. (2016) (Costa et al., 2016)	✓	✗	✗	✗	✓	✗	✓	✗	✗
Fortino et al. (2017) (Fortino et al., 2017)	✓	✗	✗	✗	✓	✓	✓	✗	✗
Josyula et al. (2017) (Josyula and Gupta, 2017)	✗	—	✓	✗	✓	✗	✓	—	✗
Savaglio et al. (2018) (Savaglio et al., 2018)	—	—	✓	✗	✓	✗	✓	—	✗
Casadei et al. (2019) (Casadei et al., 2019)	✓	✗	—	✗	✓	✗	✓	—	✗
Guerrero-Ulloa et al. (2023) (Guerrero-Ulloa et al., 2023)	✓	✗	✓	—	—	✗	—	✗	✗
Muñoz et al. (2025) (Muñoz et al., 2025)	✓	—	✗	✗	—	✗	—	—	✗
This work (2026)	—	✓	✓	✓	—	—	✓	✓	✓

infrastructure), users (business-logic-level entities with defined roles), administrators (operational management entities), and communication infrastructure (data exchange enablers).

Definition 1 (Cybersecurity in an IoT Ecosystem). *Intrinsic properties of the design, development, structure, integration, operation, and physical conditions that determine the vulnerability, robustness, resilience, and response to disturbances or malicious interventions in IoT ecosystems.*

We likewise adopt the reconceptualization of the Root of Trust (RoT) as a multi-phase *process*, rather than a static hardware component, established in Campos-Sanchez et al. (2026a): entities establish, exchange, and securely store identity parameters within a trusted environment before any operational communication begins. This process accommodates heterogeneous device capabilities—from constrained microcontrollers to cloud servers—supports both symmetric and asymmetric cryptographic schemes, and establishes the trusted context required for provisioning, authentication, and secure session establishment across all entity types. When Phase 1 (parameter generation) is instantiated through a Physically Unclonable Function (PUF) rather than random generation followed by persistent storage, Phase 2 (secure stor-

age) is correspondingly trivialized: the entity’s private identity parameter is never stored at rest—only non-secret helper data required to reconstruct it deterministically—removing the secure-storage attack surface rather than merely hardening it. Stage 6.1.3 (Section 4) instantiates this process concretely for the case study’s entity types.

4. Proposed Methodology

The methodology is organized into four phases and seven stages (Figure 2). Each stage produces documented outputs that serve as inputs to the subsequent stage, and is itself decomposed into numbered sub-processes (e.g., 2.2.4, 6.1.3) that the case-study evidence cites throughout this article; the full numbered specification is provided in the Supplementary Material (Tables S1–S4).




Figure 2: Organization of the proposed methodology: four phases (Organizational, Architecture, Implementation, Operation) comprising seven numbered stages (S1–S7). Iterative development cycles (ITER) advance Technology Readiness Level. Maintenance points M1 (implementation), M2 (deployment), and M3 (operation) trigger targeted stage re-entry. The Root of Trust is executed at Stage 6.

4.1. Phase A: Organizational (Stage 1)

This phase addresses organizational considerations for IoT adoption. Stage 1 establishes the problem to be solved, operating policies, development policies, and cybersecurity criteria. The key processes are: problem definition with cost-benefit analysis (1.1); operating policies covering functional

objectives, financial constraints, and regulatory compliance (1.2); development policies defining an iterative strategy using the TRL framework, timelines, and resource allocation (1.3); and cybersecurity policies covering criticality assessment, technological sovereignty, data sovereignty, supply chain security, and applicable government guidelines including Secure-by-Design (UK Government Security, 2024) and Privacy-by-Design (European Union Agency for Cybersecurity, 2017) principles (1.4). Processes 1.4.2–1.4.3 operationalize the organizational-strategy layer of the sovereignty framework established in Campos-Sanchez et al. (2026a), under which technological and data sovereignty are treated as emergent properties of who controls the entities’ Root of Trust process rather than as external policy overlays.

4.2. Phase B: Architecture (Stages 2–3)

Stage 2 defines the IoT ecosystem architecture at the entity level: integration diagram, communication topology, entity hierarchy, functional processes, data flow, and scalability strategy. Critically, Stage 2 instantiates the formal identity model established in Campos-Sanchez et al. (2026a), under which each entity X in ecosystem E possesses an identity:

$$I_X^E = P_{X\text{ pub}}^E \cup P_{X\text{ priv}}^E \quad (1)$$

where public parameters $P_{X\text{ pub}}^E$ serve as identifiers and private parameters $P_{X\text{ priv}}^E$ enable cryptographic operations, and authentication parameters $P_{X\text{ auth}}^E \subset I_X^E$ and shared secrets $P_{X\text{ ss}}^E \subset P_{X\text{ auth}}^E$ are derived from this identity foundation (Figure 3). This methodology extends that model with an explicit cardinality count for each entity type X : the public-parameter and private-parameter counts are:

$$n_X^E = |P_{X\text{ pub}}^E| \quad (2)$$

$$k_X^E = |P_{X\text{ priv}}^E| \quad (3)$$

so that the total identity parameter count is $n_X^E + k_X^E$ and the table columns report the two component counts and their sum.

Stage 3 specifies the services required for entity operation: secure storage subsystems, access control mechanisms, identity parameter storage strategies, and regulatory compliance integration. For constrained devices, Stage 3 requires an explicit cost/risk evaluation of persistent secure-storage options (hardware security modules, encrypted flash partitions) against physically-derived alternatives that avoid persisting a secret altogether; a Physically




Figure 3: Identity parameter set hierarchy for entity X in ecosystem E . Shared secrets P_{SS} enable cryptographic operations; authentication parameters P_{auth} prove identity; private parameters P_{priv} are never disclosed; public parameters P_{pub} serve as identifiers. In the SMG, the node’s private identity parameter (its PUF response) lives in P_{priv} but is never persisted to storage—only the non-secret mask/ECC helper data used to reconstruct it lives in ESP32 NVS, as P_{pub} .

Unclonable Function (PUF) that regenerates the private identity parameter on demand from non-secret helper data is one such alternative, and its selection at Stage 4 is a direct consequence of this Stage 3 evaluation rather than an independent design choice. Servers and users retain conventional secrets management and MFA, respectively.

4.3. Phase C: Implementation (Stages 4–5)

Stage 4 selects enabling technologies aligned with policies from Stage 1 and architecture from Stages 2–3. Technology choices are evaluated against maturity level (TRL), vulnerability landscape, and supply chain risk. For the SMG, Stage 4 selects `esp32_puflib` (Stanicek, 2022), an SRAM-based PUF implementation for the ESP32 family, as the identity-generation mechanism motivated by the Stage 3 secure-storage evaluation above. Because `esp32_puflib` is implemented as a native C/ESP-IDF component, Stage 4 also selects a C/ESP-IDF firmware architecture for the node, so that the PUF library is linked directly without a language-runtime wrapper; the same decision selects ESP-IDF’s native `mbedtls/esp_aes` for AES-256-GCM and ECDSA support. This platform choice is itself a direct, traceable consequence of the Stage 3 evaluation, not an independent implementation preference. Stage 5 implements all ecosystem elements: servers, devices, databases, communication infrastructure, IoT services, and applications. Security is enforced through secure coding practices, unit testing of security mechanisms, and validation against applicable security guidelines (Secure-by-Design principles (UK Government Security, 2024), NIST IoT Cybersecurity Guidance (Fagan et al., 2021), and OWASP IoT Top 10).

4.4. Phase D: Operation (Stages 6–7)

Stage 6 covers deployment, configuration, and testing: individual component validation, infrastructure deployment, *Root of Trust execution* for all entities (Process 6.1.3), general operation verification, and security testing (penetration testing at higher TRLs; tamper verification at TRL 3). Stage 7 addresses ongoing operation: monitoring, element maintenance, information lifecycle management, threat intelligence, element updates, incident response, and network infrastructure oversight.

4.5. Iterative Development and Maintenance

The methodology supports iterative advancement through TRL-guided cycles (ITER) and three structured maintenance points: M1 (during implementation, triggers technology re-evaluation), M2 (during deployment,

triggers code revision), and M3 (during operation, triggers architectural redesign). M1 fires when an implementation-phase review reveals that the chosen technology cannot satisfy a stage requirement; M2 fires when deployment-phase testing reveals that the implementation requires code-level revision to meet a security or functional requirement; M3 fires when operation-phase monitoring reveals that the deployed architecture cannot accommodate an observed operational pattern. This structure supports progressive maturation from proof-of-concept to production deployment.

4.6. Root of Trust Establishment

Building on the four-phase Root of Trust process of [Campos-Sanchez et al. \(2026a\)](#), we instantiate its parameter generation, secure storage, exchange, and peer-management phases as concrete protocol message sequences for both cryptographic schemes.

PUF-seeded device identity (resource-constrained devices): Rather than generating a random device key pair and persisting the private key in flash, the device’s identity key pair is derived on demand from a Physically Unclonable Function (PUF), instantiated via `esp32_puflib` ([Stanicek, 2022](#)). *Enrollment* (performed once, Table 2) reads the ESP32’s RTC FAST SRAM power-up state across repeated measurements and derives two artifacts—a stable-bit mask and error-correction (helper) data—that are stored on the device and allow the same physical response to be reconstructed, with correction, on every subsequent boot. Reconstructing the response from these artifacts alone, without the chip, is intended to be infeasible, consistent with the design goal of SRAM-based PUFs ([Guajardo et al., 2007](#)); however, repetition-code helper data is known to leak some response information ([Dodis et al., 2004](#); [Maes, 2013](#); [Delvaux et al., 2015](#)), so this is a design property, not a proven guarantee for this construction (see Limitations). On every boot, the reconstructed response deterministically derives an elliptic-curve (P-256) key pair via HKDF-based key derivation (RFC 5869) followed by rejection sampling to a uniform scalar; the device registers its public key with the server once, at enrollment, and the corresponding private key is never transmitted, stored, or otherwise disclosed—it exists only transiently in RAM for the duration of a single boot. *Session identification* (Table 3) is asymmetric on the device side and symmetric on the server side: the device signs a server-issued nonce with its private key (ECDSA), verified against its registered public key; the server proves possession of its own pinned, per-node verifier by hashing it with a device-issued nonce, checked against the

value pinned at enrollment. No proof is reused across sessions, since each is bound to a fresh nonce. The device’s NVS holds only the non-secret stable-bit mask/ECC data (P_{pub}) and the pinned per-node server verifier V_B ; its own private key (sk_A , derived fresh from PUF_A each boot) exists only transiently in RAM (Section 3).

Server identity: the server is not a good physical candidate for PUF-based identity (no comparable unclonable substrate is assumed for the RPi4), so its identity parameter is a conventionally-provisioned verifier V_B , generated fresh per enrolled device and pinned on that device as its server-authentication trust anchor. V_B is deliberately not shared across the fleet: recovering one device’s pinned V_B lets an attacker impersonate the server only to that single device. Server authentication is symmetric: possession of V_B is exactly what the server must prove per session, via the same nonce-bound hash structure that gives the device’s proof its freshness guarantee; its confidentiality on the device side is bounded by physical security (Information Disclosure, Table 6; Limitations). Device authentication, by contrast, is asymmetric: the server stores only each device’s public key, so a full server compromise discloses no value sufficient to impersonate any device (Section 5.6). The two entities’ parameters are thus deliberately asymmetric, bounding the single point of failure a fully-symmetric scheme would place on the server.

Asymmetric scheme (capable devices): Entities generate public/private key pairs, exchange public keys with identifiers, and store the peer’s public key for subsequent authentication using ISO/IEC 9798-3 mutual authentication (ISO/IEC, 2019) (Table 4).

Implementation note: the implementation source (`smg_puf_rot_node/`, device side; `smg_primary_server/rot_challenge_service.py`, server side; companion code repository, Materials Availability) uses `shs_b/V_B` for V_B , `challenge_a/challenge_b` for N_A/N_B , and `server_proof/device_proof` for R_B/R_A (`device_proof` carries a DER-encoded ECDSA signature, not a fixed-length hash); the four-message sequence and the Key_{ab} derivation above match the implementation exactly, including concatenation order and the `utils.Prehashed` convention the server uses to verify a signature computed over an already-hashed message, verified directly against the source on 2026-10-06.

Table 2: Root of Trust process: PUF-seeded device identity establishment — Phase 1 (enrollment and registration, performed once, over an isolated/physical provisioning channel). PUF_A denotes the device’s PUF response, reconstructed from RTC FAST SRAM via `esp32_puflib`; $\text{mask}_A/\text{ecc}_A$ are the non-secret stable-bit mask and error-correction helper data. (sk_A, pk_A) is the device’s P-256 key pair, derived deterministically from PUF_A ; only pk_A (public) is ever registered or stored off-device. V_B is a conventionally-provisioned, per-node server verifier (Section 4, “Server identity”), not PUF-derived and not shared across nodes.

Device	Message	Server
$\{ID_A\}$ enroll_puf() $\rightarrow \{\text{mask}_A, \text{ecc}_A\}$ (stored, public) get_puf_response() $\rightarrow \text{PUF}_A$; $(sk_A, pk_A) \leftarrow$ HKDF-derive-keypair(PUF_A)	$\{ID_A, pk_A\} \rightarrow$ $\leftarrow \{ID_B, URL/IP_B,$ $Port_B, V_B\}$	$\{ID_B, URL/IP_B, Port_B\}$ (fleet-wide, non-secret) generate fresh per-node V_B ; register $\{ID_A, pk_A, V_B\}$ store $\{ID_A, pk_A\}$ as device reference (pk_A public, not a secret)
pin (store): $\{ID_B, V_B\}$ as server reference		

Table 3: Root of Trust process: PUF-seeded device identity establishment — Phase 2 (per-boot reconstruction and mutual challenge; continues Table 2). Device authentication is asymmetric: sk_A is never transmitted after being derived, and the server verifies R_A against the device’s registered public key pk_A alone. Server authentication remains a nonce-bound hash-reconstruction proof over the per-node verifier V_B . No secret is persisted on the device between Phase 1 and Phase 2; PUF_A (and hence sk_A) is regenerated each boot.

Device	Message	Server
generate nonce N_A get_puf_response() using $\{\text{mask}_A, \text{ecc}_A\}$ $\rightarrow PUF_A$; $(sk_A, pk_A) \leftarrow$ HKDF-derive-keypair(PUF_A) verify $R_B \stackrel{?}{=} \text{SHA-256}(V_B \parallel N_A)$ using pinned V_B (server authentication); $R_A \leftarrow$ ECDSA-sign($sk_A, \text{SHA-256}(N_B \parallel N_A)$)	$\{ID_A, N_A\} \rightarrow$ $\leftarrow \{ID_B, N_B, R_B = \text{SHA-256}(V_B \parallel N_A)\}$ $\{R_A\} \rightarrow$ $\leftarrow \{\text{session_ack}\}$ (or rejection if verification failed)	generate nonce N_B compute R_B from this node’s stored V_B ECDSA-verify($pk_A, \text{SHA-256}(N_B \parallel N_A), R_A$) using registered pk_A (device identification) issue session_ack only if R_A verifies $\text{SHA-256}(V_B \parallel N_A \parallel N_B) \rightarrow Key_{ab}$ Key_{ab} held in RAM; pk_A reference retained in server storage (public, not sensitive)
$\text{SHA-256}(V_B \parallel N_A \parallel N_B) \rightarrow Key_{ab}$ Key_{ab} held only in RAM for session duration; sk_A discarded with PUF_A at end of boot		

Table 4: Root of Trust process: asymmetric key establishment.

Device	Message	Server
$\{ID_A, MAC_A, PubKey_A, PrivKey_A\}$	$\{ID_A, MAC_A, PubKey_A\}$	$\{ID_B, URL/IP_B, Port_B, PubKey_B, PrivKey_B\}$
	$\rightarrow$	
	$\leftarrow \{ID_B, URL/IP_B, Port_B, PubKey_B\}$	
secure store: $\{PubKey_B, ID_B, URL/IP_B, Port_B\}$		secure store: $\{PubKey_A, ID_A, MAC_A\}$

4.7. Threat Coverage by Design: STRIDE Analysis

We analyze how the methodology’s stages address the STRIDE (Spoofing, Tampering, Repudiation, Information Disclosure, Denial of Service, and Elevation of Privilege) threat categories (Shostack, 2014) for centralized IoT ecosystems. Unlike framework-specific threat models tied to particular technologies, this analysis maps threats to methodology stages and technology-agnostic mechanism categories, ensuring applicability across implementations.

Tables 5 and 6 map each STRIDE threat to the methodology stage(s) providing countermeasures, the mechanism category mandated by those stages, and the instantiation in the SMG case study. The adversary $\mathcal{A}$ is assumed to have network access, commodity computational resources, and knowledge of the system architecture (Kerckhoffs’s principle); goals include unauthorized data access, service disruption, entity impersonation, and credential exfiltration.

5. Case Study: DC Smart Microgrid IoT Ecosystem

5.1. Case Study Design

This case study follows a **confirmatory design** (Yin, 2014): the methodology constitutes the theoretical proposition; the SMG implementation constitutes the empirical test. Evidence is collected from four sources following Yin’s principle of multiple evidence sources (Table 7).

Table 5: STRIDE threat model: mapping to methodology stages and SMG instantiation — Part 1 (Spoofing, Tampering, Repudiation).

Threat	Attack Vector	Methodology Stage(s)	Mechanism Category	SMG Instantiation	Residual Risk
Spoofing	Entity impersonation; forged credentials	2.2.2 (auth mech.), 2.2.3 (RoT), 6.1.3 (RoT execution)	Cryptographic mutual authentication using entity-specific, PUF-derived identity commitments	PUF-derived identity (<code>esp32_puflib</code>) + asymmetric device proof/symmetric server proof; no persistent secret in NVS; handshake before any data accepted; confirmed in 30/30 real trials (Section 6.3)	Low (requires physical device compromise)
Tampering	Data modification in transit; sensor reading manipulation	2.2.4 (secure comms), 5.2.1 (secure dev.)	Authenticated encryption with integrity binding per message; random nonces	AES-256-GCM per message; random 12-byte nonce; GCM tag detects single-message modification	Low per-message (GCM tag); Medium for reordering/replay (no message-sequence chain, see Limitations)
Repudiation	Denial of sent data; dispute of control commands	2.2.2 (session audit), 3.1 (data service), 7.1 (monitoring)	Session-bound data exchange with server-side session logging	Server logs each session under <code>node_id</code> and timestamp (Stage 7.1); evidentiary, not an independently-verifiable chain (trade-off, see Limitations)	Medium (logs can be disputed without independent cryptographic chain verification)

Table 6: STRIDE threat model (continued from Table 5) — Part 2 (Information Disclosure, Denial of Service, Elevation of Privilege).

Threat	Attack Vector	Methodology Stage(s)	Mechanism Category	SMG Instantiation	Residual Risk
Information Disclosure	Key extraction; credential exfiltration; traffic interception	1.4.2–1.4.3 (sovereignty), 2.2.1 (identity params), 5.2.1 (secure dev.)	Data sovereignty policy; no persistent device-own secret to disclose; application-layer encryption	Node identity reconstructed at boot from the PUF, never persisted; only non-secret helper data and the pinned server verifier V_B in NVS; all data AES-256-GCM encrypted	Node identity: Low (server stores only pk_A , insufficient to forge a signature). Server impersonation from a captured unit: Medium (V_B readable from its NVS; see Limitations)
Denial of Service	Authentication flooding; bus load disruption; node isolation	7.1 (monitoring), 7.2.1 (anomaly detect.), 2.1.8 (hierarchy)	Anomaly-based trust scoring; degraded-mode operation; entity isolation policy	Trust score decay on repeated failures; UNTRUSTED nodes isolated (TCP dropped); last PWM duty maintained on timeout (original server modules; not yet ported to the RoT server, see Limitations)	Medium (physical-layer DoS outside software scope)
Elevation of Privilege	Node masquerading as admin; unauthorized config access	2.1.8 (entity hierarchy), 3.2.1 (access control), 2.2.2 (auth)	Entity hierarchy with API-enforced access; token type includes entity class	Node session tokens permit only POST sensor data; admin endpoints reject non-admin tokens	Low (token type embedded in authenticated session)

Table 7: Data collection protocol: evidence sources.

Source	Evidence Type	Content
Design documents	Documentary	Stage outputs: policies, architecture diagrams, identity tables, service specifications
Firmware artifacts	Physical artifact	ESP32 C/ESP-IDF modules: sensor reading, EMS, identity (PUF enrollment/reconstruction, <code>esp32_puflib</code>), communication (AES-256-GCM), metrics
Execution logs	Archival records	Authentication logs, RoT deployment logs, functional test reports, <code>/methodology.csv</code> output
Measurements	Quantitative	Handshake latency, sampling rate, payload size, heap stability, OWASP coverage

5.2. Research Questions

5.3. Case Context

Application domain: The SMG manages distributed DC electrical energy generation, storage, and consumption across two co-located nodes in a controlled laboratory environment. The growing integration of renewable energy sources (solar panels, batteries) into local grids requires intelligent, automated management to balance generation and demand in real time without relying on centralized national grid infrastructure. This application is classified as critical infrastructure (Methodology Stage 1, Process 1.4.1), requiring the highest cybersecurity tier.

Physical architecture: The SMG consists of two secondary control agent nodes and one primary agent (Raspberry Pi 4B). Each secondary node connects a renewable generation element to a shared 12 V DC bus with local energy storage (Table 9). The primary agent runs on a Raspberry Pi 4B connected to the DC bus as a pure load via a DC/AC inverter (12 V DC $\rightarrow$ 120 V AC) followed by a USB-C AC/DC adapter.

Digital architecture: The ecosystem uses a centralized communication topology (Stage 2.1.4) with fog processing (Stage 2.1.7): each ESP32

Table 8: Case study research questions.

ID	Question	Addressed by
CSQ1	Does the methodology provide sufficient guidance to architect a complete, secure IoT ecosystem for a real industrial application?	Stage-by-stage traceability; completeness of outputs; functional test results
CSQ2	Does the parameterized identity model enable unambiguous security mechanism definition across heterogeneous entity types?	Identity parameter table; authentication implementation; coverage across 4 entity types
CSQ3	Does the Root of Trust Process establish verifiable secure communication between ecosystem entities?	RoT execution evidence; handshake latency measurements; mutual challenge-response verification
CSQ4	What is the methodology’s coverage of known IoT security threats relative to established standards?	OWASP IoT Top 10 mapping; security mechanism traceability

Table 9: Secondary node hardware components.

Component	Role	Specification
ESP32 NodeMCU	IoT + EMS controller	Dual-core 240 MHz, 520 KB SRAM, 4 MB Flash, built-in WiFi; powered by local battery via DC-DC buck (12V→5V)
Epcom AGM VRLA	Local energy storage + ESP32 supply	PL-18-12, 12 V, 18 Ah; also supplies ESP32 via buck converter
IBFIA DC-DC SSR ×3	Energy flow control	3–32 VDC control, 5–60 VDC load, 40 A; driven by ESP32 LEDC on GPIO27/26/25 at 1 kHz, 10-bit
Schneider MCB ×3	Overcurrent protection	2-pole, one per SSR channel
Solar charge controller	Energy flow control	Commercial MPPT controller (handles droop control and battery-charging voltage regulation autonomously)
ACS712-05B ×2	Current sensing	Hall-effect, 185 mV/A, 16-sample median oversampling
Resistive divider ×2	Voltage sensing	30 k Ω /7.5 k Ω (5:1 ratio); reads up to 16.5 V on 3.3 V ADC

executes local sensor fusion and a 9-rule fuzzy Energy Management System (EMS), while the RPi4 runs global coordination with a 27-rule security-aware EMS. The entity hierarchy is: IoT Nodes (lowest privilege) < Users < Administrators. Development target: TRL 3 (proof of concept with physical elements).

5.4. Stage-by-Stage Traceability

Tables 10 and 11 document the application of each methodology stage to the SMG, showing the decision made at each stage and the measurable artifact produced. This traceability table constitutes the primary evidence for CSQ1.

5.5. Identity Parameter Definition (Stage 2.2.1)

Applying Equations (1)–(3) to the SMG entities produces a 15-parameter identity model (Table 12). The public/private/authentication parameter separation directly determined the authentication mechanism (PUF-derived commitment using the private, non-persistent parameter) and the secure storage requirement for the IoT Node: none, beyond the non-secret mask/ECC helper data already classified as public.

The entity-specific separation of identity parameters produced a direct, unambiguous mapping to implementation artifacts: `node_id`, `node_uuid`, `maskA`, and `eccA` correspond to distinct NVS keys written at enrollment time; no parameter redefinition or additional design decisions were required at the specification level. In the present case study, P_{auth} and P_{priv} coincide for every entity type (all private parameters participate in authentication, as shown in Table 12); this is a case-study-specific coincidence, not a general property of the model—Campos-Sanchez et al. (2026a) define only the broader relation $P_{auth} \subset I$ (Figure 3), so future case studies may include a private parameter not used for authentication (e.g., a long-term audit key), in which case $P_{auth} \subsetneq P_{priv}$ rather than $P_{auth} = P_{priv}$.

[†]**Note on the PUF response as a non-persistent private parameter:** the formal identity model $I_X^E = P_{X_{pub}}^E \cup P_{X_{priv}}^E$ (Section 4) classifies a parameter as private or public by its confidentiality role, not by whether it is stored. The IoT Node’s PUF response satisfies P_{priv} ’s role—it is never disclosed and underlies the device’s authentication proof—while never existing as data at rest: it is regenerated from `maskA`/`eccA` (both P_{pub}) on every boot. This is the concrete instance of the Phase 2 (secure storage) trivialization noted in Section 3.

Table 10: Stage-by-stage traceability: methodology process $\rightarrow$ SMG decision $\rightarrow$ artifact — Part 1 (Stages 1–3).

Stage	Methodology Process	Decision Made in SMG	Artifact
1 Technological Paradigm	1.1 Establish problem	Manage distributed renewable energy nodes securely in a 12 V DC microgrid	Problem specification
	1.1.2 Entities involved	5 entity types: IoT nodes, server, users, admins, DC bus	Entity list
	1.4.1 Criticality	Classified as critical infrastructure (electrical energy)	High-security cybersecurity policy
	1.4.2 Technological sovereignty	Self-hosted; no third-party cloud; ESP32 + local RPi4 server	Architecture constraint record
<i>Of the seven conceptual-framework element types (Section 3), the SMG instantiates five: storage/communication infrastructure are subsumed into the server, and IoT services into the device/server software. Of the five entity types, four (nodes, server, users, admins) carry digital identity (Table 12); the DC bus is tracked via P3 sensor-data integrity instead.</i>			
2 Digital Ecosystem	2.1.4 Communication topology	Centralized: all nodes report to primary agent	Architecture diagram
	2.1.7 Processing topology	Fog: local EMS on ESP32; global coordination on RPi4	Processing topology record
	2.2.1 Identity definition	4 entity types, 15 identity parameters total (non-uniform per type; Table 12)	Identity parameter table
	2.2.2 Authentication mechanism	Asymmetric (ECDSA) device proof + symmetric per-node server proof, both PUF-seeded; AES-256-GCM session encryption	Authentication diagram
	2.2.3 Root of Trust process	4-phase RoT designed for node-server pairs	RoT execution log
3 IoT Services	3.1 Service definition	4 services: data ingestion, EMS coordination, auth API, monitoring	Service specification
	3.2.1 Access control	Token-based; nodes: POST only; users: GET; admins: full	Access control matrix

Table 11: Stage-by-stage traceability (continued from Table 10) — Part 2 (Stages 4–7).

Stage	Methodology Process	Decision Made in SMG	Artifact
4 Tech. Modular Integration	4.1.1 TRL planning	TRL 3 target for ITER 1; lab. proof of concept	TRL milestone record
	4.1.2 Platform selection	ESP32 + C/ESP-IDF: WiFi, SHA-256 HW, NVS, <code>esp32_puflib</code> as native component	Technology selection rationale
	4.1.3 Communication	WiFi IEEE 802.11 b/g/n + TCP; application-layer mutual-authentication protocol (<code>rot_session.c/rot_challenge_service.py</code>)	Protocol selection record
	4.2 Vulnerability mitigation	ESP-IDF v5.4.2 pinned (per the flashed firmware build log); nonce reuse fixed	Risk register
5 Applications & Devices	5.1.1 Server	REST API + time-series database + monitoring dashboard on RPi4	Server deployment
	5.1.2 Devices	ESP32 C/ESP-IDF firmware: identity (PUF), Root of Trust, secure-communication modules; sensor/EMS not yet implemented (see Limitations)	Firmware artifact
	5.2.1 Secure development	NVS stores only non-secret mask/ECC helper data (no persistent secret); <code>esp_fill_random()</code> nonces; no credentials in source	Code review record
6 Testing & Deployment	6.1.3 RoT execution	Validated end-to-end on 1/2 nodes (NODE_01; NODE_02 pending), $N = 30$ trials, 30/30 success (Section 6.3)	PUF bench (Section 6.3)
	6.2.1 Functional testing	Boot, PUF reconstruction, mutual authentication (P2), encrypted exchange (P3) all pass across $N = 30$ trials; sensor/EMS tests not yet applicable	PUF bench (Section 6.3)
	6.2.2 Tamper verification	NVS tampering confirmed as auth. failure (single negative test; full pen. testing deferred to ITER 2)	Security verification log
7 Operation & Mgt.	7.1.1 Monitoring	Real-time V, I, P, PWM duty per node; authentication status; session history	Dashboard artifact
	7.2.2 Update elements	OTA update policy defined with authentication requirement (ITER 2 target)	Update policy

Table 12: Identity parameter table: Stage 2.2.1 output for the SMG. $\text{mask}_A/\text{ecc}_A$: non-secret stable-bit mask and error-correction helper data generated at enrollment by `esp32_puflib`, stored in ESP32 NVS. PUF response[†]: the device’s physically-derived private identity parameter, reconstructed from $\text{mask}_A/\text{ecc}_A$ on every boot and never persisted (see note below). $V_B^\dagger$: the server’s per-node authentication verifier (see note below). `node_uuid` is public: it identifies the device in Phase 2 messages but is not itself part of the cryptographic authentication proof (Table 3). Following Eqs. (2)–(3), $n = n_X^E$ is the public-parameter count, $k = k_X^E$ the private-parameter count, and $\text{Total} = n_X^E + k_X^E$ the total identity-parameter count.

Entity Type	Class	P_{pub}	P_{priv}	P_{auth}	n	k	Total
IoT Node (ESP32)	Device	<code>node_id</code> , <code>mask_A</code> , <code>ecc_A</code> , <code>node_uuid</code>	PUF response [†]	PUF response [†]	4	1	5
IoT Server (RPi4)	Server	<code>url/ip</code> , <code>port</code> , <code>id_b</code>	$V_B^\dagger$	$V_B^\dagger$	3	1	4
User	User	<code>username</code>	<code>user_uuid</code> , <code>password</code>	<code>user_uuid</code> , <code>password</code>	1	2	3
Administrator	Admin	<code>username</code>	<code>admin_uuid</code> , <code>password</code>	<code>admin_uuid</code> , <code>password</code>	1	2	3
Total		9	6	6			15

[‡]**Note on V_B as the server’s private/authentication parameter:** the server’s row describes its own proof of identity to a device, made by demonstrating knowledge of V_B (Table 3) — the server-side analogue of the Node’s PUF response. `id_b` (fleet-wide, non-secret) sits under P_{pub} with `url/ip` and `port`. The device separately pins its own copy of V_B as a trust anchor (Section 4, “Server identity”); this is a stored reference to *another* entity’s parameter, not an additional Node parameter, so it is not double-counted.

5.6. Security Mechanism Analysis: STRIDE-to-Protocol Mapping

The methodology-level STRIDE analysis (Section 4.7) identifies mechanism categories that address each threat category at the framework level, independent of implementation technology. This subsection demonstrates how the SMG’s Stage 4 technology selection instantiates those mechanisms through three authentication and data-exchange protocol phases (P1–P3), establishing a complete traceability chain from threat identification to deployed countermeasure.

Scope note: Phases P1–P3 constitute an *implementation-specific* Stage 4 technology choice for the SMG (see Introduction scope note). Any protocol providing equivalent mechanism categories—pre-deployment identity isolation (P1), nonce-bound mutual identification without key transmission (P2), and per-message authenticated encryption (P3)—would satisfy the same STRIDE and OWASP properties against the same methodology requirements. This design intentionally omits a cryptographic session-chain/audit mechanism; Tables 5/6 and the Limitations (Section 7) record the resulting Repudiation/I8 trade-off explicitly.

Table 13 describes each phase’s security role, mechanism, and STRIDE/OWASP coverage. Table 14 summarizes the STRIDE coverage matrix across phases.

Spoofing is countered at two layers: P1 ties valid identity to physical access to the specific enrolled chip rather than a copyable stored secret; P2’s nonce-bound challenge requires fresh, per-session proof of each side’s verifier without revealing the physical response or verifier itself, confirmed bidirectionally before any session key is derived.

Tampering is detected at message creation time by the AES-256-GCM authentication tag on every message (P3); P2’s mutual identification additionally binds the session key to both parties’ proven identities. The scheme has no session-level hash chain, so retrospective re-verification of message ordering is unavailable (downgraded to $\sim$ in Table 14; see Limitations).

Table 13: SMG authentication protocol phases: security role, mechanism, and coverage mapping (P1–P3). STRIDE: S = Spoofing, T = Tampering, R = Repudiation, ID = Information Disclosure, DoS = Denial of Service, EoP = Elevation of Privilege.

Ph.	Name	Security Role	Mechanism	STRIDE	OWASP
P1	Enrollment	Identity establishment before any network exposure	<code>enroll_puf()</code> (<code>esp32_puflib</code> (Stanicek, 2022)): 20 RTC-FAST-SRAM measurements $\rightarrow$ stable-bit mask + 8 $\times$ -repetition-code ECC helper data (both non-secret), written to ESP32 NVS via USB serial on isolated workstation; <code>node_uuid</code> (UUID-4) also written (WiFi credentials are a compile-time firmware constant, not NVS-provisioned, see Limitations); device public key pk_A registered with the server, which generates and pins a fresh per-node verifier V_B in return; no raw PUF response or private key is ever written to flash	S, ID	I1, I9
P2	Root of Trust	Per-boot response reconstruction and nonce-bound mutual identification without key transmission	<code>get_puf_response()</code> reconstructs PUF_A (421 bytes; Section 6.3); $(sk_A, pk_A) \leftarrow \text{HKDF-derive-keypair}(PUF_A)$; device proves possession of sk_A via ECDSA signature over a server-issued nonce; server proves possession of this node's V_B via a nonce-bound hash (Table 3); $Key_{ab} \leftarrow \text{SHA-256}(V_B \parallel N_A \parallel N_B)$; neither the raw response, sk_A , nor V_B is ever transmitted; confirmed in 30/30 real trials, mean latency 558.0 ± 184.3 ms (Section 6.3)	S	I1, I3
P3	Data Exchange	Per-message authenticated encryption	$c_i \leftarrow \text{AES-256-GCM}(Key_{ab}, \text{nonce}_i, m_i)$ via ESP-IDF <code>mbedtls/esp_aes</code> ; a fresh random 12-byte nonce per message; the GCM authentication tag detects any single-message modification	S, T, ID	I2, I3, I6, I7

Table 14: STRIDE coverage matrix across protocol phases. ✓ = primary countermeasure (first line of defense); ~ = secondary/partial contribution; — = not applicable.

STRIDE Threat	P1	P2	P3
Spoofing	✓	✓	—
Tampering	—	~	✓
Repudiation	—	—	~
Info. Disclosure	✓	✓	✓
Denial of Service	—	~	—
Elev. of Privilege	✓	✓	—

Repudiation coverage is evidentiary rather than cryptographic: P3 messages are authenticated in transit and the server logs each session (`node_id`, timestamp, Stage 7.1), but no independently-verifiable chain binds messages to a session without trusting the server’s logs, a deliberate trade-off for a smaller on-device footprint (see Limitations).

Information Disclosure is addressed at the source (P1: no key ever exists in storage, only non-secret helper data), during derivation (P2: the physical response is reconstructed transiently in RAM, never transmitted, and neither verifier is ever sent), and in transit (P3: per-message AES-256-GCM with random nonces).

Denial of Service is partially addressed at P2: the server’s first reply (R_B) is one nonce plus one SHA-256 hash, cheap regardless of the requester’s authenticity, but no further session state is committed—and no `session_ack` issued—until the device’s ECDSA proof R_A verifies, imposing an asymmetric cost on an attacker holding a session open. Volumetric and application-layer flooding beyond the identification gate remain outside P1–P3’s scope; Stage 7 trust-based isolation (Table 6) is the mandated mechanism, implemented in the SMG’s original server modules but not yet ported to the current RoT server (see Limitations).

Elevation of Privilege is prevented at P1, since entity type is determined by which identification channel a party uses, not a field inside a shared message: nodes authenticate exclusively through the P1–P3 RoT protocol, while users and administrators use a separate, credential-based API (Table 12); no message in Tables 2/3 carries an entity-type field. P3 API access is bound

to the P2 session key, so a node session cannot invoke user- or admin-scoped endpoints.

6. Empirical Evaluation

6.1. CSQ1: Architecture Guidance Sufficiency

Traceability completeness: The stage-by-stage traceability tables (Tables 10 and 11) indicate that all 7 stages produced documented, verifiable artifacts—23 distinct process-to-artifact rows across the four phases. No stage required design decisions beyond what the methodology’s processes prescribe.

Functional validation: Table 15 reports the functional test results, drawn from the $N = 30$ real hard-reset bench trial (Section 6.3) and the dedicated tamper-verification test.

Iterative development confirmation: One maintenance moment was identified and resolved during real-hardware bench validation on SMG_NODE_01 (Section 6.3): an enrollment-state-confusion bug causing an unbounded re-enrollment loop, fixed by a code-level revision before the $N = 30$ trial reported in Table 17. This confirms the methodology’s M2 maintenance pathway operates as specified at TRL 3 (see Discussion for the M1/M2 classification and further incidents pending SMG_NODE_02 validation).

Answer to CSQ1: The methodology provided sufficient guidance to produce a complete, operational IoT ecosystem. All 7 stages produced documented outputs with no external guidance required.

6.2. CSQ2: Identity Model Unambiguity

Applying the identity model equations to 4 entity types produced 15 distinct parameters with zero overlap between entity classes (Table 12). The public/private/authentication tripartite structure directly determined: (a) what is public non-secret helper data stored in NVS (`maskA`, `eccA`, `node_id`, `node_uuid`), (b) what is never persisted and is instead reconstructed on every boot (the PUF response), and (c) what constitutes the device’s authentication proof (the reconstructed PUF response alone; `node_uuid` identifies the device but is not itself cryptographically verified). Table 16 summarizes these identity model metrics; of the SMG’s 2 nodes, 1 (SMG_NODE_01) has been provisioned and validated under this model so far, with SMG_NODE_02 pending (Section 6.3).

Table 15: Functional test results (Stage 6.2.1). FT-01–FT-04 pass rates are over the $N = 30$ hard-reset bench trial (Section 6.3); FT-05/FT-06 are single dedicated negative tests. Sensor/EMS functional tests are not yet applicable, since those modules are not implemented in the current firmware (see Limitations).

ID	Description	Result
FT-01	Boot: PUF identity reconstructed from NVS helper data (mask/ECC); no raw key or PUF response in NVS	Pass (30/30)
FT-02	WiFi connect: associates to the configured access point and obtains an IP address	Pass (30/30)
FT-03	P2 Root of Trust: asymmetric device proof (ECDSA) + symmetric server proof both verify; mutual authentication established	Pass (30/30)
FT-04	P3 session: AES-256-GCM encrypted demo message sent and acknowledged	Pass (30/30)
FT-05	Tamper verification: deliberate NVS parameter corruption consistently produces authentication failure	Pass (1/1)
FT-06	Cross-chip cloning: SMG_NODE_01’s entire NVS (PUF helper data, server reference) copied onto SMG_NODE_02’s physical chip; server rejects the resulting device proof (signature invalid)	Pass (1/1)
Overall pass rate		32/32 = 100%

Table 16: CSQ2 identity model metrics.

Metric	Value
Entity types defined	4
Total identity parameters	15
Public parameters (P_{pub})	9
Private parameters (P_{priv})	6
Authentication parameters (P_{auth})	6
Nodes provisioned via RoT	1/2
Parameter collisions (overlapping definitions)	0
Additional design decisions required at Stage 5	0

Answer to CSQ2: The identity model is sufficient and unambiguous. All 4 entity types received complete, non-overlapping identity definitions that translated directly to implementation artifacts without modification.

6.3. CSQ3: Root of Trust Verifiability

The RoT process was executed and bench-validated end-to-end on real hardware for SMG_NODE_01 (ESP32-D0WD-V3 rev 3.1), connected over the real SMG_Primary WiFi access point to the real RPi4 primary agent (`rpi4-server`) running the server-side counterpart `rot_challenge_service.py`. **Validation to date covers one of the SMG’s two nodes; SMG_NODE_02 validation remains future work.** Post-provisioning, node authentication was verified by deliberate NVS parameter tampering—corrupted parameters consistently produced authentication failures, confirming that the RoT creates a verifiable security boundary.

Table 17 reports results from $N = 30$ independent hard-reset trials against a single, fixed enrollment (the stable-bit mask and ECC helper data were generated once via `enroll_puf()` and not regenerated between trials); each trial therefore measures per-boot reconstruction-and-handshake reliability, not enrollment-to-enrollment variability, and exercises Phases P2 and P3 of the protocol (Section 5.6) against the real server on every reset, without repeating P1 enrollment.

Three findings stand out. First, the end-to-end protocol—PUF reconstruction, nonce-bound mutual identification in both directions (V_B and the device’s ECDSA signature against pk_A , in all 30 trials), and AES-256-GCM data exchange—succeeded in every trial, confirming a stable identity (and key pair) despite per-attempt measurement noise. Second, the bimodal reconstruction-attempt pattern is explained, not merely noted: the raw boot logs show attempt 1 always follows a `POWERON_RESET` (the bench harness’s external EN-pin pulse via RTS, 100 ms, USB power maintained throughout—not a true power cycle) and attempt 2 always follows a `DEEPSLEEP_RESET` (the firmware’s own internal retry). The two attempts measure the RTC FAST SRAM’s power-up transient under systematically different reset conditions, not independent draws from one noise process; why a deep-sleep wake yields a more favorable readout is a question about the underlying SRAM physics this bench does not resolve. One harmless side effect: since the firmware opens its TCP connection before attempting reconstruction, every attempt-1 failure leaves an orphaned connection the server logs as a timeout (33 such timeouts across this run). Third, one trial (run 20) measured 1530 ms

Table 17: PUF-based RoT real-hardware bench trial: SMG_NODE_01, $N = 30$ independent hard-reset trials against one fixed enrollment. Attempt-1 diagnostics are from `esp32_puflib`’s own threshold check (hamming-weight must exceed 48.5%; raw bit-error rate, over the full 4096-byte/32768-bit RTC-SRAM region before stable-bit masking, must stay below 0.15%).

Metric	Value ($N = 30$)	Notes
End-to-end success rate	30/30 (100%; 95% Clopper–Pearson lower bound 88.4%)	Mutual identification (P2) + AES-256-GCM demo message exchange (P3), confirmed on both device and server logs; repeated trials on one device/enrollment, not independent samples (see discussion below)
Reconstruction attempts per trial	2/2 in every trial (0/30 succeeded on attempt 1)	Perfectly bimodal across all 30 trials; see discussion below
Attempt-1 hamming-weight	49.915 ± 0.100 % (range 49.753–50.143 %)	Threshold > 48.5 %; passed in all 30 attempts
Attempt-1 raw bit-error rate	0.397 ± 0.042 % (range 0.317–0.479 %)	Threshold < 0.15 %; over the full 32768-bit region before stable-bit masking; failed in all 30 first attempts, passed on every second attempt
Attempt-1 error bits (of 32768 raw bits)	130.1 ± 13.6 (range 104–157)	$130.1/32768 = 0.397$ %, the same figure as the row above
Handshake latency (TCP connect $\rightarrow$ session established)	558.0 ± 184.3 ms (range 510–1530, median 520)	P2 asymmetric device proof (ECDSA) + symmetric server proof; one trial (run 20) measured 1530 ms, all other 29 fell in 510–600 ms; see discussion below
WiFi association latency	3542.7 ± 8.3 ms (range 3530–3560)	ESP-IDF stack retried the first association attempt in nearly every trial; a WiFi-stack characteristic, not a RoT protocol property
Device key-pair (pk_A) reproducibility	Consistent with a single registered value in 30/30 trials 33	All 30 trials authenticated against the one server-registered public key; confirmed indirectly, since the raw key material is not logged

against 510–600 ms for the other 29; the gap falls entirely in the first server round-trip, before ECDSA signing begins, consistent with an occasional server-side scheduling delay. All latencies are read from the ESP-IDF boot log’s millisecond timestamp (quantized to 10 ms). The mean is driven by run 20’s outlier (excluding it: 524.5 ± 16.8 ms for the remaining 29); the median (520 ms) is the more robust statistic. The 30 trials are repeated measures on one device and one enrollment, not independent samples, bounding how far these figures generalize beyond SMG_NODE_01 (see Limitations). Raw logs and the parsed summary are published in the companion artifact repository at `case_study_SMG/Implementation/smg_puf_rot_node/bench_results_20261005/` (Materials Availability).

Answer to CSQ3: The RoT process established verifiable mutual authentication for SMG_NODE_01 across $N = 30$ independent hard-reset trials (30/30 end-to-end success, Table 17): the device authenticated the server via a nonce-bound hash-reconstruction proof over a per-node verifier, and the server authenticated the device via an ECDSA signature verified against the device’s registered public key, with neither the device’s private key nor the server’s per-node verifier ever transmitted. Mean handshake latency was 558.0 ± 184.3 ms (median 520 ms), dominated by the ECDSA sign/verify step and an occasional server-side scheduling delay (see above) rather than by WiFi round-trip time. Validating SMG_NODE_02 under this scheme remains future work.

6.4. CSQ4: Threat Coverage

Tables 18 and 19 map each OWASP IoT Top 10 item to the methodology stage(s) mandating the countermeasure, the protocol phase(s) that provide the implementation-level evidence (cross-referenced with the phase analysis in Section 5.6 and Table 13), and the specific SMG mechanism.

Answer to CSQ4: Unweighted OWASP IoT Top 10 coverage of 7/10 items fully addressed and 3/10 partially addressed. Six of the seven fully-covered items (I1–I3, I6–I7, I9) have direct evidence from protocol phases P1–P3; the remaining one (I5: component vetting) is addressed by Stage 4 policy artifacts. The three partial items are I4 (secure update: policy defined, signed-OTA deferred to ITER 2), I8 (device management: server-side logging without a verifiable chain, a trade-off for a smaller footprint—see Limitations), and I10 (physical hardening: outside a software methodology’s scope). The coverage score is an unweighted count, not a severity-weighted metric (B12, Section 6.5).

Table 18: OWASP IoT Top 10 coverage analysis — Part 1 (I1–I5). Stage = methodology process mandating the countermeasure. Phase(s) = protocol phase(s) providing implementation evidence (see Table 13); — = addressed by policy/deployment, not protocol.

#	OWASP IoT Risk	Stage(s)	Phase(s)	Evidence in SMG	Coverage
I1	Weak/hardcoded passwords	2.2.1	P1, P2	No password or key is used for device identity: reconstructed at boot from a PUF using only non-secret mask/ECC helper data, never a stored password; confirmed working in 30/30 real trials (Section 6.3). The separate WiFi credential is a compile-time firmware constant, set per build rather than a vendor default, but not NVS-provisioned (see Limitations)	Full
I2	Insecure network services	2.2.4	P2, P3	P2 establishes mutual identification and a session key before any data is accepted; P3 encrypts every message with per-message random 12-byte nonce (AES-256-GCM); no plaintext channel exists	Full
I3	Insecure ecosystem interfaces	3.2.1	P2	P2 nonce-bound mutual identification, both parties verified before any session key is derived; node vs. user/admin access separated by identification channel (STRIDE EoP, Section 5.6); trust scoring exists but not yet ported to the RoT server (Limitations)	Full
I4	Lack of secure update	7.2.2	—	OTA update policy (Stage 7.2.2): admin authentication + firmware signature required; implementation scoped to ITER 2 (policy artifact only, no signed-OTA deployment at TRL 3)	Partial
I5	Insecure/outdated components	4.2	—	ESP-IDF v5.4.2 pinned; no deprecated crypto primitives (ECDSA/SHA-256; AES-256-GCM via <code>mbedtls</code>)	Full

Table 19: OWASP IoT Top 10 coverage analysis (continued from Table 18) — Part 2 (I6–I10).

#	OWASP IoT Risk	Stage(s)	Phase(s)	Evidence in SMG	Coverage
I6	Insufficient privacy	1.4.3 + 2.2.4	P3	Data sovereignty policy (Stage 1.4.3); self-hosted RPi4; no PII; sensor data only; P3 encrypts all telemetry in transit with authenticated encryption	Full
I7	Insecure data transfer/storage	3.1.1 + 2.2.4	P3	P3 provides AES-256-GCM encryption in transit; no device-identity secret at rest (PUF-derived). V_B is unencrypted in NVS; the WiFi credential is an unencrypted firmware constant (standard ESP32 practice; flash encryption is an ITER 2 candidate for both)	Full
I8	Lack of device management	7.1	—	Stage 7.1 server monitors heap, send rate, and session status from server-side logs (node_id, timestamps); without a cryptographic chain, this is an operational log rather than an independently verifiable activity audit	Partial
I9	Insecure default settings	2.2.3 + 6.1.3	P1	No WiFi/server comms. until enrollment data exists in NVS; identity proof derives from the chip’s physical PUF response, not a stored/default value; confirmed by cloning SMG_NODE_01’s entire NVS onto SMG_NODE_02’s physical chip, which the server correctly rejected (FT-06)	Full
I10	Lack of physical hardening	—	— 36	Out of scope for software methodology; addressed at org. level (Stage 1.4.1); JTAG disable and secure boot specified for TRL 5+	Partial
Coverage score (un-weighted)					7/10 full + 3/10 partial

6.5. Expert Validation

A structured questionnaire (14 Likert items, B1–B14, 5-point scale from 1 = strongly disagree to 5 = strongly agree, plus 3 open questions) was distributed to domain experts to supplement the case study evidence with independent professional judgment. Fifteen evaluators completed the questionnaire; fourteen consented to their responses being used as validation evidence (one declined consent and was excluded prior to analysis), exceeding the upper bound of the planned 6–10 evaluator target. Six requested anonymity and are referenced anonymously below. Evaluators held primary expertise in cybersecurity, IoT systems design/development, smart-grid/energy management, software engineering, or embedded systems, with 1–15 years of professional or research experience and academic backgrounds ranging from B.Sc./B.Eng. to Ph.D., from institutions in Mexico, Spain, Portugal, and Canada. Eight of the fourteen are affiliated with institutions unconnected to the author team (CINVESTAV, ESCOM-IPN, a national Instituto Tecnológico, ECOSUR, Universidad Iberoamericana, Universidade de Aveiro, University of New Brunswick), making the panel majority-independent; the remaining six are affiliated with the author institutions (INAOE, home institution; CEIT-BRTA, co-author affiliation).

Quantitative results: the overall mean across all 196 Likert responses (14 evaluators $\times$ 14 items) was 3.98/5, with item means ranging from 3.57 (B12: OWASP coverage sufficiency) and 3.71 (B1: overall structure sufficiency) to 4.43 (B10: hash-chain integrity guarantees) and 4.36 (B3: stage-by-stage artifact verifiability; B14: more systematic integration than existing methodologies). No item fell below the questionnaire’s neutral midpoint (3.0) on average; among items whose mechanism matches the current implementation, B8 (dynamic RoT execution is a verifiable identity mechanism, 4.00/5) corresponds directly to CSQ3. The independent-evaluator subset (N=8) returned an overall mean of 3.89/5, modestly below the affiliated subset (4.11/5) and the full panel; the direction of this gap indicates mild affiliation skew, which the majority-independent composition of the panel mitigates.

Qualitative findings that qualify the results: B12 (OWASP coverage sufficiency) returned the lowest item mean (3.57/5): two evaluators noted the 9/10 OWASP score (as presented) is an unweighted count, not severity-weighted, and does not establish that I10 (physical hardening) is the least critical omission; subsequent reclassification of I4 and I8 to partial yields the revised 7/10 full + 3/10 partial reported in CSQ4 (Section 7). B1

Table 20: Expert questionnaire Likert results (N=14; 1 = strongly disagree, 5 = strongly agree).

Item	Statement (abbreviated)	Mean
B1	Four-phase, seven-stage structure gives sufficient guidance without external references	3.71
B2	Organizational phase (Stage 1) adds practical value ahead of architecture/technology decisions	3.79
B3	Stage-by-stage artifact requirements make correct methodology application verifiable	4.36
B4	TRL-guided iterative model with M1–M3 maintenance points is practical	3.86
B5	Entity-based identity model is clearer than role/policy-based approaches	3.86
B6	12-parameter identity model generalizes beyond the SMG case study	3.86
B7	RoT-as-process abstraction accommodates heterogeneous device capabilities	4.29
B8	Dynamic RoT execution (Process 6.1.3) is a verifiable identity mechanism	4.00
B9	Six-phase authentication/data-exchange protocol (P1–P6) is logically complete	3.86
B10	Hash-chain mechanism (P4/P6) gives meaningful, feasible integrity guarantees	4.43
B11	Technology-independent STRIDE mapping is more useful than technology-specific models	4.07
B12	90% OWASP IoT Top 10 coverage is sufficient security-by-design evidence at TRL 3	3.57
B13	Parameterized, technology-agnostic design applies beyond energy management	3.79
B14	More systematic security integration than existing IoT methodologies	4.36
Overall mean (196 responses)		3.98

Note (B12): item text reflects OWASP coverage as presented to evaluators (9/10 full + 1/10 partial = 90%); the coverage assessment reported in CSQ4 (Section 6) is 7/10 full + 3/10 partial, since I4 (secure-update policy) and I8 (device management) are both classified as partial rather than full.

Note (B6, B7, B9, B10): these item texts, as presented to evaluators, describe a 12-parameter identity model and a six-phase P1–P6 protocol that includes a hash-chain integrity mechanism. The current implementation (Section 6.3) uses a 15-parameter model, phases P1–P3, and no hash chain. The panel was not re-surveyed against the current implementation, so these means should be read as validating the design choices as presented, not the current mechanisms.

(structure sufficiency) returned the second-lowest mean (3.71/5): three evaluators scored it 2/5, noting the methodology specifies *what* each stage must produce but not *how*—per-artifact templates or acceptance criteria would close this gap, echoed by several open-question (C3) responses. One evaluator qualified B6, noting the parameter count is an artifact of the SMG’s four entity types rather than a fixed model limit (14 parameters, Table 12), with generalization to other entity types untested. Two evaluators gave B13 (cross-domain generalizability) dissenting scores (2/5, 1/5), treating it as an unvalidated hypothesis (reflected in External Validity and Limitations below). One evaluator questioned whether same-team case-study evidence is sufficient without an independently-managed development team; the SMG implementation was carried out by multiple co-authors in distinct roles (Author Contributions), which partially addresses this. Two evaluators outside cybersecurity specifically (software-engineering methodology and general IoT systems) noted several protocol/threat-model items fell outside their expertise; their ratings are retained but should be weighted accordingly. The panel’s majority-independent composition and subset means are reported above; the residual affiliation skew, convenience-sampling constraint, and preliminary-data caveat are acknowledged as limitations below.

7. Discussion

This work addresses two barriers to secure IoT ecosystem development: the security activation gap at the operational level, and the absence of a structured development methodology that integrates cybersecurity from the earliest organizational stage. The results of the confirmatory case study and expert validation are discussed below in relation to each research question.

CSQ1 — Methodology guidance: The stage-by-stage traceability (Tables 10 and 11) demonstrates that the methodology produces actionable decision points at every stage. The organizational phase (Stage 1) carries the highest leverage: classifying the SMG as critical infrastructure (Process 1.4.1) directly constrained the authentication scheme (requiring mutual cryptographic authentication), the storage approach (requiring isolation of identity material from firmware, met at TRL 3 by the NVS-partition/PUF design; a dedicated secure element is an ITER 2 requirement, see Limitations), and the access control granularity (requiring node, user, and admin distinction). Absent an explicit organizational stage, these constraints might

not have been systematically propagated to architectural and implementation decisions.

CSQ2 — Identity model: The formal separation of identity into P_{pub} , P_{priv} , and P_{auth} subsets yielded a technology-agnostic security specification that translated without modification to ESP32 NVS storage for public helper data and to SHA-256 hashing (via mbedtls) for the private, non-persistent PUF response. The 15-parameter model produced zero ambiguity across 4 entity types, confirming the model’s expressiveness for heterogeneous IoT ecosystems, including resource-constrained devices.

CSQ3 — Root of Trust verifiability: The P2+P3 mutual authentication handshake completed in 558.0 ± 184.3 ms (median 520 ms) across $N = 30$ real hard-reset trials on SMG_NODE_01, well within the firmware-guaranteed 1 s maximum inter-send interval, with 30/30 trials succeeding end-to-end (Table 17). The handshake latency is dominated by the ECDSA sign/verify step together with an occasional server-side scheduling delay (one trial, run 20, measured 1530 ms against 510–600 ms for the other 29; see Section 6.3), rather than by WiFi round-trip time. Send success rate was 100%, and the device’s derived key pair authenticated correctly against the server’s registered public key in every trial, confirming a stable identity despite per-attempt PUF measurement noise.

CSQ4 — Threat coverage: The unweighted OWASP IoT Top 10 coverage (7/10 full; 3/10 partial: I4 secure-update policy only, I8 device management without a verifiable chain, I10 physical hardening) indicates the methodology systematically addresses known IoT threat categories. The STRIDE analysis operates at two levels: the methodology-level mapping (Tables 5/6, Section 4.7) shows every threat addressed by a mechanism mandated by at least one stage; the protocol-level mapping (Tables 13–14, Section 5.6) traces each mechanism to its implementing phase. All six categories are addressed, four by multiple independent phases. The three partial OWASP items reflect two deliberate trade-offs (I4, I8—see Limitations) and one inherent scope boundary (I10); the coverage figure is unweighted, not severity-weighted.

Maintenance cycles: One maintenance moment was triggered during Stages 5–6 (an M2 firing for code-level revision, Section 6.3), confirming the methodology’s M2 pathway operates as specified at TRL 3. The issue was caught during deployment-phase bench testing, before the reported trial; the case study does not establish a counterfactual about whether it would have surfaced only later absent that structure, since bench testing of this

kind is standard practice regardless. M1 and M3 were not triggered and remain empirically unvalidated, available for future maintenance points and the operation-phase TRL transition.

7.1. Limitations

- **Single topology:** Only centralized topology was implemented at TRL 3. Decentralized and mesh topologies supported by Stage 2.1.4 remain to be validated empirically.
- **TRL 3 scope:** The proof of concept operates in a controlled laboratory environment. Environmental variation, scalability beyond 2 nodes, and production-grade hardening (JTAG disable, secure boot) are deferred to ITER 2.
- **Primary-agent power dependency:** In the TRL 3 bench configuration the RPi4 primary agent is powered from the same DC bus it coordinates (via a DC/AC inverter); SSR duty transients and node resets cause primary-agent brownouts that can terminate sustained bench trials early. This coupling constrains sustained concurrent operation and is an artifact of the laboratory power arrangement rather than the methodology; ITER 2 will use an independent primary-agent supply.
- **Forward secrecy (per-node symmetric server verifier):** The RoT’s session-key derivation, $Key_{ab} = \text{SHA-256}(V_B \parallel N_A \parallel N_B)$, depends on the per-node server verifier V_B and the handshake nonces exchanged in the clear; it does not depend on the device’s private key sk_A . Consequently, an attacker who later recovers this node’s V_B from server storage, and has recorded a past handshake’s nonces, can recompute that session’s key, even though the device’s own identity remains protected (Section 5.6). This is a property of the per-node symmetric half of the RoT at TRL 3; deriving Key_{ab} from an ephemeral ECDH exchange instead of V_B is the ITER 2 remedy.
- **Device-side PUF entropy:** The server stores only pk_A and V_B (Section 4), neither of which recreates sk_A if disclosed; however, sk_A ’s security still depends on the actual (not merely designed) entropy of the PUF response it derives from. An enrollment on one SMG node measured a 421-byte (~ 3368 -bit) stable response, somewhat above the library’s ~ 380 – 400 -byte reference-hardware estimate (chip-to-chip variation, discussed below). Raw length is not min-entropy: repetition-code helper data is known

to leak some response information (Dodis et al., 2004; Maes, 2013; Delvaux et al., 2015), so 3368 raw bits do not imply 256 bits of adversarial min-entropy **[TODO: min-entropy assessment pending]**: a dedicated assessment (e.g., NIST SP 800-90B) remains the direct way to confirm sufficient min-entropy for the 256-bit scalar it seeds.

- PUF library build and bench status (validated on one of two nodes):** `esp32_puflib` (Stanicek, 2022) is a native C/ESP-IDF component and is linked directly into the node’s C/ESP-IDF firmware (`smg_puf_rot_node`), without a language-runtime wrapper. The identity/Root-of-Trust/secure-communication firmware and its RPi4 counterpart (`rot_challenge_service.py`) build cleanly, and the firmware has been flashed to, enrolled on, and bench-validated end-to-end against one of the SMG’s two physical ESP32 units: $N = 30$ real hard-reset trials against the real RPi4 server reached 30/30 end-to-end success (Section 6.3). `SMG_NODE_02` has not yet been enrolled under this scheme in its own right, so two-node reconstruction-reliability validation (Table 17-style bench statistics for a second chip) remains open. Physical distinctness across chips, however, was directly confirmed by cloning `SMG_NODE_01`’s entire NVS onto `SMG_NODE_02`’s hardware: applying `SMG_NODE_01`’s helper data to `SMG_NODE_02`’s physical RTC-SRAM failed the library’s reliability threshold on the first attempt (29.9% raw error rate against a 0.15% requirement) and, even after a passing retry, produced a key pair the server correctly rejected (FT-06, `c4_negative_test/` in the artifact repository). The current firmware also uses a compile-time placeholder node identifier shared by all units; it identifies the node on the wire but does not enter the authentication proof, and is to be replaced by per-unit assignment before multi-node deployment.
- Sensor/EMS modules not yet implemented (C/ESP-IDF firmware scope):** The `smg_puf_rot_node` C/ESP-IDF firmware currently implements only the identity (PUF), Root of Trust, and secure-communication layers; fuzzy Energy Management System and sensor-acquisition logic have not yet been implemented in this firmware. A full end-to-end SMG deployment requires this additional work, distinct from and in addition to the hardware bench validation above. The same gap applies to the server side: the primary agent’s anomaly-based trust scoring and quarantine logic (Stage 7.2.1) exists in `smg_primary_server`’s original module set (`auth_service.py`, `ems_coordinator.py`) but has not yet been ported to

`rot_challenge_service.py`, the server counterpart of the current RoT protocol; DoS/I3 coverage claims elsewhere in this paper that cite trust scoring describe the methodology-mandated mechanism and its prior instantiation, not a property of the currently-deployed RoT server.

- **No cryptographic session-chain audit (deliberate design simplification):** The SMG’s Root of Trust does not implement a per-message hash chain to provide a retrospectively-verifiable, third-party-auditable record of session integrity. Message-level integrity is provided by the AES-256-GCM authentication tag (P3), but there is no cryptographic chain binding messages within a session, and Repudiation/I8 (device management) coverage is therefore evidentiary (server-side session logs) rather than independently verifiable (Table 5, Table 19). Introducing a lightweight, chain-free sequencing mechanism (e.g., a monotonic per-session counter bound into the GCM associated data) without the overhead of a full chain audit is a candidate ITER 2 refinement if reviewers or subsequent evaluation identify this trade-off as unacceptable.
- **PUF reliability on the deployed hardware (reconstruction reliability measured on one node; thermal/second-node characterization pending):** `esp32_puflib`’s enrollment procedure targets a design bit-flip probability of 0.001–0.005 via a stable-bit mask and 8×-repetition-code error correction, yielding an estimated ~380–400-byte stable response in the library’s own design notes for its reference hardware. Independent enrollment runs on `SMG_NODE_01` have produced stable responses of 418–421 bytes — consistently reproducible, but somewhat above that figure, a plausible chip-to-chip variation rather than a conformance failure, since the library documents its figure as an estimate, not a hard specification. Reconstruction reliability was measured over an $N = 30$ hard-reset bench round (Table 17, Section 6.3); the resulting bimodal attempt pattern is explained at the reset-condition level, not the circuit level, in that section’s discussion. What remains unmeasured is (a) sensitivity to temperature and supply-voltage variation, and (b) whether the same pattern and response length hold on `SMG_NODE_02` and across independent enrollments **[TODO: temperature/voltage and cross-node data pending]**: this remains pending before the reliability claim extends beyond `SMG_NODE_01` under laboratory-ambient conditions.
- **Single application domain:** Energy management is the only domain

tested. The expert questionnaire (N=14, Section 6.5) supports the technology-agnostic design intent (B13 mean 3.79/5) but also recorded two dissenting scores on this same item (2/5 and 1/5), with the dissenting evaluators characterizing cross-domain generalizability as an unvalidated hypothesis; empirical validation in additional domains remains future work.

- **Physical hardening:** I10 (physical hardening) and hardware-based Root of Trust (e.g., ATECC608 secure element) are not addressed at TRL 3 but are specified as ITER 2 requirements.
- **Identity lifecycle (methodology gap):** The methodology as presented does not specify a complete identity lifecycle process (parameter rotation, revocation, recovery after compromise without full re-provisioning, ownership transfer, and decommissioning). This is a gap in the methodology itself—not solely a TRL 3 scope limitation—since recovery from a compromised node currently requires re-running the full Root of Trust. An expert evaluator (Section 6.5) identified this as the single addition with the greatest potential impact on practical applicability; it is the highest-priority extension for the next methodology revision and is deferred to ITER 2.
- **User/administrator password hashing:** Users and administrators authenticate via a plain SHA-256 hash of their password (Table 12), not a salted, slow key-derivation function (bcrypt, scrypt, Argon2) — weaker than current practice, and a candidate ITER 2 hardening item; it does not affect the PUF-based IoT node scheme, which uses no password.
- **WiFi credential provisioning:** The node’s WiFi SSID/password are compile-time constants in the firmware source, set per build rather than read from NVS; unlike the PUF-derived identity material, they are not entity-specific or rotatable without reflashing. Moving them into NVS alongside the enrollment data is a candidate ITER 2 hardening item.
- **Expert validation:** Fourteen evaluators completed the questionnaire (14 Likert items, 3 open questions), exceeding the planned 6–10 target, yielding an overall agreement mean of 3.98/5 (Section 6.5). Eight of the fourteen are affiliated with institutions unconnected to the author team (majority-independent), though recruitment remained convenience-sampled; the independent subset (3.89/5) rated modestly below the affiliated subset (4.11/5), a mild skew the panel composition mitigates. At least

two affiliated evaluators (Acknowledgements) are MSc students in the authors’ own department, a power relationship beyond simple affiliation this study does not separately quantify. The condensed reference document distributed to evaluators (June 24, 2026) predates the final bench trials and used more conservative figures. One additional respondent declined consent and was excluded.

7.2. Threats to Validity

Table 21: Threats to validity.

Type	Specific Threat	Mitigation
Internal validity	SMG decisions may reflect author preference rather than methodology guidance	Stage-by-stage traceability (Tables 10 and 11) documents explicit methodology process → decision chain for each output
External validity	Single domain (energy) + single topology (centralized) limits generalizability	Methodology is technology-agnostic by design; expert questionnaire (N=14, mean 3.98/5, Section 6.5) supports this design intent, though three evaluators noted that generalization beyond the validated domain and entity set remains an untested hypothesis
Construct validity	TRL 3 does not demonstrate production readiness	TRL 3 is the stated objective of ITER 1; the methodology explicitly supports iterative maturation toward TRL 5+
Reliability	Physical prototype measurements subject to environmental variation	$N = 30$ real hard-reset trials on SMG_NODE_01 (Section 6.3): handshake latency std 184.3 ms (one scheduling-delay outlier of 30), end-to-end success 30/30; single-node scope and ambient laboratory conditions are disclosed limitations

8. Conclusions

This work has presented a security-by-design methodology for centralized IoT ecosystems that addresses the security activation gap by integrating cybersecurity systematically from the earliest organizational stage. The methodology’s four phases and seven stages produce documented, traceable artifacts; the entity-centric identity model provides technology-agnostic security specifications applicable to heterogeneous devices; and the dynamic

Root of Trust process establishes verifiable cryptographic identity before any operational communication begins.

Empirical validation through a confirmatory case study (Yin, 2014) applying the methodology to a DC Smart Microgrid IoT ecosystem yields four findings. First, all seven stages produced complete, traceable outputs with zero external guidance required (CSQ1). Second, the 15-parameter identity model across four entity types produced zero ambiguity and translated directly to implementation artifacts (CSQ2). Third, the RoT process established verifiable mutual authentication: a PUF-based scheme (`esp32_puflib` (Stanicek, 2022)) that reconstructs each node’s identity from non-persistent physical characteristics rather than a stored key, combined with an asymmetric (ECDSA) device-authentication proof and a per-node symmetric server-authentication proof, independently validated end-to-end on real hardware for one of the two nodes across $N = 30$ hard-reset trials (30/30 success, mean handshake latency 558.0 ± 184.3 ms); the scheme’s bimodal reconstruction-attempt pattern is left as an open question, and validating SMG_NODE_02 remains future work (CSQ3, Section 6.3). Fourth, the methodology achieved unweighted OWASP IoT Top 10 coverage of 7/10 items fully addressed and 3/10 partially addressed (I4 secure-update policy only; I8 device management, a trade-off for a smaller on-device footprint; I10 physical hardening outside software scope) (CSQ4). This case study evidence is supplemented by an expert questionnaire (N=14, Section 6.5) that returned an overall agreement mean of 3.98/5, with the strongest agreement on methodology completeness and systematic security integration (B3, B14 means 4.36/5) and qualified agreement on cross-domain generalizability and threat-coverage weighting.

The methodology’s parameterized, technology-agnostic structure is designed to support application across diverse IoT domains; cross-domain generalization is a design intent supported by expert opinion (B13 mean 3.79/5, with dissenting 2/5 and 1/5 ratings characterizing it as an unvalidated hypothesis) rather than a demonstrated property, and empirical validation in additional domains remains future work. Future work will extend validation to: ITER 2 (TRL 5+ with hardware security elements and signed-OTA implementation), cross-domain deployment (healthcare, industrial automation), an expanded expert evaluator pool with improved institutional independence, formal cryptographic protocol verification of the SMG’s PUF-based authentication scheme, resolution of the identity-lifecycle gap (rotation, revocation, recovery), evaluation of per-node verifier wrapping or HSM-

backed server storage to reduce central-server single-point-of-failure risk, and a lightweight chain-free sequencing mechanism to recover stronger Repudiation/I8 coverage without the on-device footprint of a full cryptographic session-chain audit.

Declarations

Materials Availability

The ESP32 firmware (`smg_puf_rot_node/`) and primary agent software (`smg_primary_server/`) are available under Apache 2.0 license in the companion code repository (Campos-Sanchez et al., 2026b): <https://github.com/RazielCS/smg-secure-microgrid>. The complete research artifacts — manuscript source, case-study documentation, raw bench-trial measurements (canonical CSVs and serial logs under `case_study_SMG/Implementation/smg_puf_rot_node/`), and the expert-validation instruments — are available in the companion artifact repository (Campos-Sanchez et al., 2026c): <https://github.com/RazielCS/smg-secure-microgrid-paper-artifacts>. Device- and server-side secrets (the real PUF verifiers, NVS images, and WiFi credentials) are excluded from both repositories and replaced with sanitized `.example.json` templates and placeholder values; see each repository’s README for the exact exclusion list.

Data Statement

The raw data underlying this work’s quantitative findings — bench-trial CSVs, serial logs, packet captures, and the expert-validation instruments (anonymized where required by evaluator consent; see Acknowledgements) — have been deposited in, and are cited and linked from, the companion repositories above (Campos-Sanchez et al., 2026b,c). No part of the underlying data is withheld from these deposits other than the device/server secrets noted under Materials Availability, which are excluded for security rather than confidentiality reasons and are not needed to reproduce or verify any reported result.

Competing Interests

The authors declare no competing financial interests or personal relationships that could have appeared to influence this work.

Funding

R. Campos-Sanchez thanks SECIHTI for scholarship support (CVU 623274).

Declaration of Generative AI and AI-Assisted Technologies in the Manuscript Preparation Process

During the preparation of this work, the authors used Claude (Anthropic, Claude Code) as an AI-assisted software engineering and writing-support tool for: implementing, debugging, and bench-testing the case study’s firmware and server software; statistical analysis of the real-hardware bench-trial data reported in Section 6.3; drafting and revising manuscript text, including cross-section consistency verification; and preparing the companion code and artifact repositories cited in Materials Availability. An independent AI coding assistant (OpenCode, running third-party models) was additionally used to perform an independent critical-review pass of the manuscript, whose findings the authors evaluated and selectively incorporated. After using these tools, the authors reviewed, verified, and edited all AI-assisted output and take full responsibility for the content of the published article. AI tools are not listed as authors or co-authors of this work.

Acknowledgements

The authors thank the following experts for completing the validation questionnaire (Section 6.5) and consenting to attribution: Jose de Jesus Zapata Lara (CINVESTAV); Jorge Andres Larracoechea, PhD (Asociacion Centro Tecnologico Ceit); Claudia García-Blanquel, PhD (ESCOM-IPN); Edwin Pacheco, MSc student in Security Sciences and Technologies (INAOE); Agustin Alberto Ahumada Ramirez, MSc student in Security Sciences and Technologies (INAOE); Yair Romero Lopez, M.Eng. (INAOE); Alfonso Martínez-Cruz, PhD (SECIHTI-INAOE); and Eder Toledo, MSc (ECOSUR). Six additional evaluators contributed responses under conditions of anonymity and are not named here per their request. One further respondent declined consent to the use of their responses and was excluded prior to analysis.

Author Contributions

R. Campos-Sanchez: Conceptualization, investigation, formal analysis, writing — original draft, implementation. **L. Hernandez-Martinez:** Supervision, validation, writing — review & editing. **C. Feregrino-Uribe:** Supervision, validation, writing — review & editing. **A. Penuelas-Angulo:**

Conceptualization, formal analysis, writing — original draft. **D. Cruz-Aguilar:** Conceptualization, validation, implementation, writing — review & editing.

Vitae

[TODO: author biographies and photographs pending]

A short (maximum 100-word) biography and a passport-type photograph for each of the 5 authors — R. Campos-Sanchez, L. Hernandez-Martinez, C. Feregrino-Uribe, A. Penuelas-Angulo, D. Cruz-Aguilar — are to be supplied by each author before submission; photographs are to be submitted as separate figure files, not embedded here.

References

- Abraham, D., Houmb, S.H., Erdodi, L., 2025. Cyber-attacks on energy infrastructure—a literature overview and perspectives on the current situation. *Applied Sciences* 15, 9233.
- Al-Fuqaha, A., Guizani, M., Mohammadi, M., Aledhari, M., Ayyash, M., 2015. Internet of things: A survey on enabling technologies, protocols, and applications. *IEEE Communications Surveys & Tutorials* 17, 2347–2376.
- Antonakakis, M., April, T., Bailey, M., Bernhard, M., Bursztein, E., Cochran, J., Durumeric, Z., Halderman, J.A., Invernizzi, L., Kallitsis, M., Kumar, D., Lever, C., Ma, Z., Mason, J., Menscher, D., Seaman, C., Sullivan, N., Thomas, K., Zhou, Y., 2017. Understanding the Mirai botnet, in: *Proceedings of the 26th USENIX Security Symposium*, USENIX Association, Vancouver, BC, Canada. pp. 1093–1110.
- Ayala, I., Amor, M., Fuentes, L., Troya, J.M., 2015. A software product line process to develop agents for the IoT. *Sensors* 15, 15640–15660.
- Bakiris, E., Oikonomakis, G., 2025. IoT-driven security solutions, in: *Cybersecurity: 9th European Interdisciplinary Cybersecurity Conference, EICC 2025*, Rennes, France, June 18–19, 2025, *Proceedings*, Springer Nature. p. 275.

- Blanco, C., Rosado, D.G., Varela-Vaca, Á.J., Gómez-López, M.T., Fernández-Medina, E., 2023. Onto-CARMEN: Ontology-driven approach for Cyber-Physical System security requirements meta-modelling and reasoning. *Internet of Things* 24, 100989.
- Campos-Sanchez, R., Feregrino-Urbe, C., Hernandez-Martinez, L., Penuelas-Angulo, A., 2026a. An identity-centric paradigm for the internet of things: Security, trust, and sovereignty by design. *IoT Submitted for publication*.
- Campos-Sanchez, R., Hernandez-Martinez, L., Feregrino-Urbe, C., Penuelas-Angulo, A., Cruz-Aguilar, D., 2026b. smg-secure-microgrid: Companion code repository for “closing the security activation gap”. Software repository, <https://github.com/RazielCS/smg-secure-microgrid>.
- Campos-Sanchez, R., Hernandez-Martinez, L., Feregrino-Urbe, C., Penuelas-Angulo, A., Cruz-Aguilar, D., 2026c. smg-secure-microgrid-paper-artifacts: Research artifacts and bench-trial data for “closing the security activation gap”. Dataset repository, <https://github.com/RazielCS/smg-secure-microgrid-paper-artifacts>.
- Casadei, R., Tsigkanos, C., Viroli, M., Dustdar, S., 2019. Engineering resilient collaborative edge-enabled IoT, in: 2019 IEEE International Conference on Services Computing (SCC), IEEE. pp. 36–45.
- Cicirelli, F., Fortino, G., Guerrieri, A., Spezzano, G., Vinci, A., 2017. Meta-modeling of smart environments: from design to implementation. *Advanced Engineering Informatics* 33, 274–284.
- Cloud Security Alliance, 2023. IoT Security Controls Framework, Version 3. Technical Report. Cloud Security Alliance. URL: <https://cloudsecurityalliance.org/research/working-groups/internet-of-things>.
- Costa, B., Pires, P.F., Delicato, F.C., 2016. Modeling IoT applications with SysML4IoT, in: 2016 42nd Euromicro Conference on Software Engineering and Advanced Applications (SEAA), IEEE. pp. 157–164.
- Delvaux, J., Gu, D., Schellekens, D., Verbauwheide, I., 2015. Helper data algorithms for PUF-based key generation: Overview and analysis. *IEEE*

- Transactions on Computer-Aided Design of Integrated Circuits and Systems 34, 889–902.
- Dodis, Y., Reyzin, L., Smith, A., 2004. Fuzzy extractors: How to generate strong keys from biometrics and other noisy data. *Advances in Cryptology (EUROCRYPT 2004)*, 523–540.
- European Telecommunications Standards Institute, 2020. ETSI EN 303 645 V2.1.1: Cyber security for consumer internet of things: Baseline requirements. European Standard. URL: https://www.etsi.org/deliver/etsi_en/303600_303699/303645/02.01.01_60/en_303645v020101p.pdf.
- European Union Agency for Cybersecurity, 2017. Baseline Security Recommendations for IoT in the Context of Critical Information Infrastructures. Technical Report. European Union Agency for Cybersecurity (ENISA).
- Fagan, M., Marron, J., Brady, K., Cuthill, B., Megas, K., Herold, R., 2021. IoT Device Cybersecurity Guidance for the Federal Government: Establishing IoT Device Cybersecurity Requirements. NIST Special Publication 800-213. National Institute of Standards and Technology (NIST). Gaithersburg, MD. doi:[10.6028/NIST.SP.800-213](https://doi.org/10.6028/NIST.SP.800-213).
- Fernandez, E.B., Brazhuk, A., 2024. A critical analysis of Zero Trust Architecture (ZTA). *Computer Standards & Interfaces* 89, 103832.
- Fernando, Y., Tseng, M.L., Wahyuni-Td, I.S., de Sousa Jabbour, A.B.L., Chiappetta Jabbour, C.J., Foropon, C., 2023. Cyber supply chain risk management and performance in industry 4.0 era: information system security practices in Malaysia. *Journal of Industrial and Production Engineering* 40, 102–116.
- Fortino, G., Russo, W., Savaglio, C., Shen, W., Zhou, M., 2017. Agent-oriented cooperative smart objects: From IoT system design to implementation. *IEEE Transactions on Systems, Man, and Cybernetics: Systems* 48, 1939–1956.
- Fortino, G., Savaglio, C., Spezzano, G., Zhou, M., 2020. Internet of things as system of systems: A review of methodologies, frameworks, platforms, and tools. *IEEE Transactions on Systems, Man, and Cybernetics: Systems* 51, 223–236.

- Guajardo, J., Kumar, S.S., Schrijen, G.J., Tuyls, P., 2007. FPGA intrinsic PUFs and their use for IP protection, in: *Cryptographic Hardware and Embedded Systems (CHES 2007)*, Springer. pp. 63–80.
- Guerrero-Ulloa, G., Rodriguez-Dominguez, C., Hornos, M.J., 2023. Agile methodologies applied to the development of Internet of Things (IoT)-based systems: A review. *Sensors* 23, 790.
- ISO/IEC, 2019. ISO/IEC 9798-3:2019 information technology — security techniques — entity authentication — part 3: Mechanisms using digital signature techniques. International Organization for Standardization. 5th edition.
- Javanmardi, S., Shojafar, M., Mohammadi, R., Alazab, M., Caruso, A.M., 2023. An SDN perspective IoT-fog security: A survey. *Computer Networks* 229, 109732.
- Josyula, S.K., Gupta, D., 2017. A new security methodology for internet of things, in: *2017 International Conference on Computing, Communication and Automation (ICCCA)*, IEEE. pp. 613–618.
- Licata, J., McWhite, R., Calloway, L., Gilbert, D., Anderson, M., Snyder, J., Miller, J., 2025. Developing Security, Privacy, and Cybersecurity Supply Chain Risk Management Plans for Systems. Technical Report. National Institute of Standards and Technology.
- Maes, R., 2013. *Physically Unclonable Functions: Constructions, Properties and Applications*. Springer.
- Makhdoom, I., Abolhasan, M., Lipman, J., Liu, R.P., Ni, W., 2018. Anatomy of threats to the internet of things. *IEEE Communications Surveys & Tutorials* 21, 1636–1675.
- Manate, B., Forti, F., Moore, P., 2014. Applying the Prometheus methodology for an Internet of Things architecture, in: *2014 IEEE/ACM 7th International Conference on Utility and Cloud Computing*, IEEE. pp. 435–442.
- Merzouk, S., Cherkaoui, A., Marzak, A., Nawal, S., 2020. IoT methodologies: comparative study. *Procedia Computer Science* 175, 585–590.

- Moreta, N., Aragon, D., Ona, S., Jaramillo, A., Ibarra, J., Jahankhani, H., 2023. Comparison of cybersecurity methodologies for the implementing of a secure IoT architecture, in: *Cybersecurity in the Age of Smart Societies: Proceedings of the 14th International Conference on Global Security, Safety and Sustainability*, London, September 2022, Springer. pp. 9–29.
- Morin, B., Harrand, N., Fleurey, F., 2017. Model-based software engineering to tame the IoT jungle. *IEEE Software* 34, 30–36.
- Muñoz, M., Torres, M., Gil, J.D., Guzmán, J.L., 2025. An Internet of Things platform for heterogeneous data integration: Methodology and application examples. *Journal of Network and Computer Applications* 240, 104197.
- Okika, N., Nwatuzie, G.A., Odozor, L., Oni, O., Idoko, I.P., 2025. Addressing IoT-driven cybersecurity risks in critical infrastructure to safeguard public utilities and prevent large-scale service disruptions. *International Journal of Innovative Science and Research Technology* 10.
- Patel, P., Cassou, D., 2015. Enabling high-level application development for the Internet of Things. *Journal of Systems and Software* 103, 62–84.
- Ray, R.K., Chowdhury, F.R., Hasan, M.R., 2024. Blockchain applications in retail cybersecurity: Enhancing supply chain integrity, secure transactions, and data protection. *Journal of Business and Management Studies* 6, 206–214.
- Saqib, M., Moon, A.H., 2023. A systematic security assessment and review of Internet of Things in the context of authentication. *Computers & Security* 125, 103053.
- Sassalou, E.N., Vasileiadis, S., Kazazis, S.A., Protogerou, G., Varvitsiotis, N., Karas, D.S., Giannetsos, T., Tsintzos, S., 2025. A PUF-based root-of-trust for resource-constrained IoT devices, in: *2025 IEEE International Conference on Cyber Security and Resilience (CSR)*, IEEE. pp. 824–831.
- Savaglio, C., Fortino, G., Gravina, R., Russo, W., 2018. A methodology for integrating internet of things platforms, in: *2018 IEEE International Conference on Cloud Engineering (IC2E)*, IEEE. pp. 317–322.
- Shahinzadeh, G., Shahinzadeh, H., Tanwar, S., 2024. Security and privacy issues in the Internet of Things: A comprehensive survey of protocols,

- standards, and the revolutionary role of blockchain, in: 2024 8th International Conference on Smart Cities, Internet of Things and Applications (SCIoT), IEEE. pp. 59–67.
- Shostack, A., 2014. Threat Modeling: Designing for Security. Wiley, Indianapolis, IN.
- Sorri, K., Mustafee, N., Seppanen, M., 2022. Revisiting IoT definitions: A framework towards comprehensive use. *Technological Forecasting and Social Change* 179, 121623.
- Spanoudakis, N., Moraitis, P., 2015. Engineering ambient intelligence systems using agent technology. *IEEE Intelligent Systems* 30, 60–67.
- Stanicek, O., 2022. esp32_puflib: An SRAM-based physical unclonable function library for the ESP32 microcontroller family. Software library, ESP-IDF component, https://github.com/staniond/esp32_puflib.
- Turner, S.W., Karakus, M., Guler, E., Uludag, S., 2023. A promising integration of SDN and blockchain for IoT networks: A survey. *IEEE Access* .
- UK Government Security, 2024. Secure by design principles. <https://www.security.gov.uk/guidance/secure-by-design/principles/>. Accessed: 30 April 2024.
- Visintin, A., 2024. Secure IoT and the Identity of Things. Ph.D. thesis. Università degli Studi di Padova.
- Yin, R.K., 2014. Case Study Research: Design and Methods. 5th ed., SAGE Publications, Thousand Oaks, CA.
- Zanasi, C., Russo, S., Colajanni, M., 2024. Flexible zero trust architecture for the cybersecurity of industrial IoT infrastructures. *Ad Hoc Networks* , 103414.
- Zhou, I., Makhdoom, I., Shariati, N., Raza, M.A., Keshavarz, R., Lipman, J., Abolhasan, M., Jamalipour, A., 2021. Internet of things 2.0: Concepts, applications, and future directions. *IEEE Access* 9, 70961–71012.